

Sizing LLM inference systems

How many GPUs do you need for inference?



About Us



Dmitry Mironov, EMEA

- Senior Deep Learning Solutions Architect @ NVIDIA - Supporting deployment of AI / Deep Learning solutions
- Focusing on large scale efficient deployment and inference



Sergio Perez, EMEA

- Senior Deep Learning Solutions Architect @ NVIDIA - Supporting delivery of AI / Deep Learning solutions
- Focusing on quantization in training and inference



Agenda

Sizing for Inference can get a bit complicated

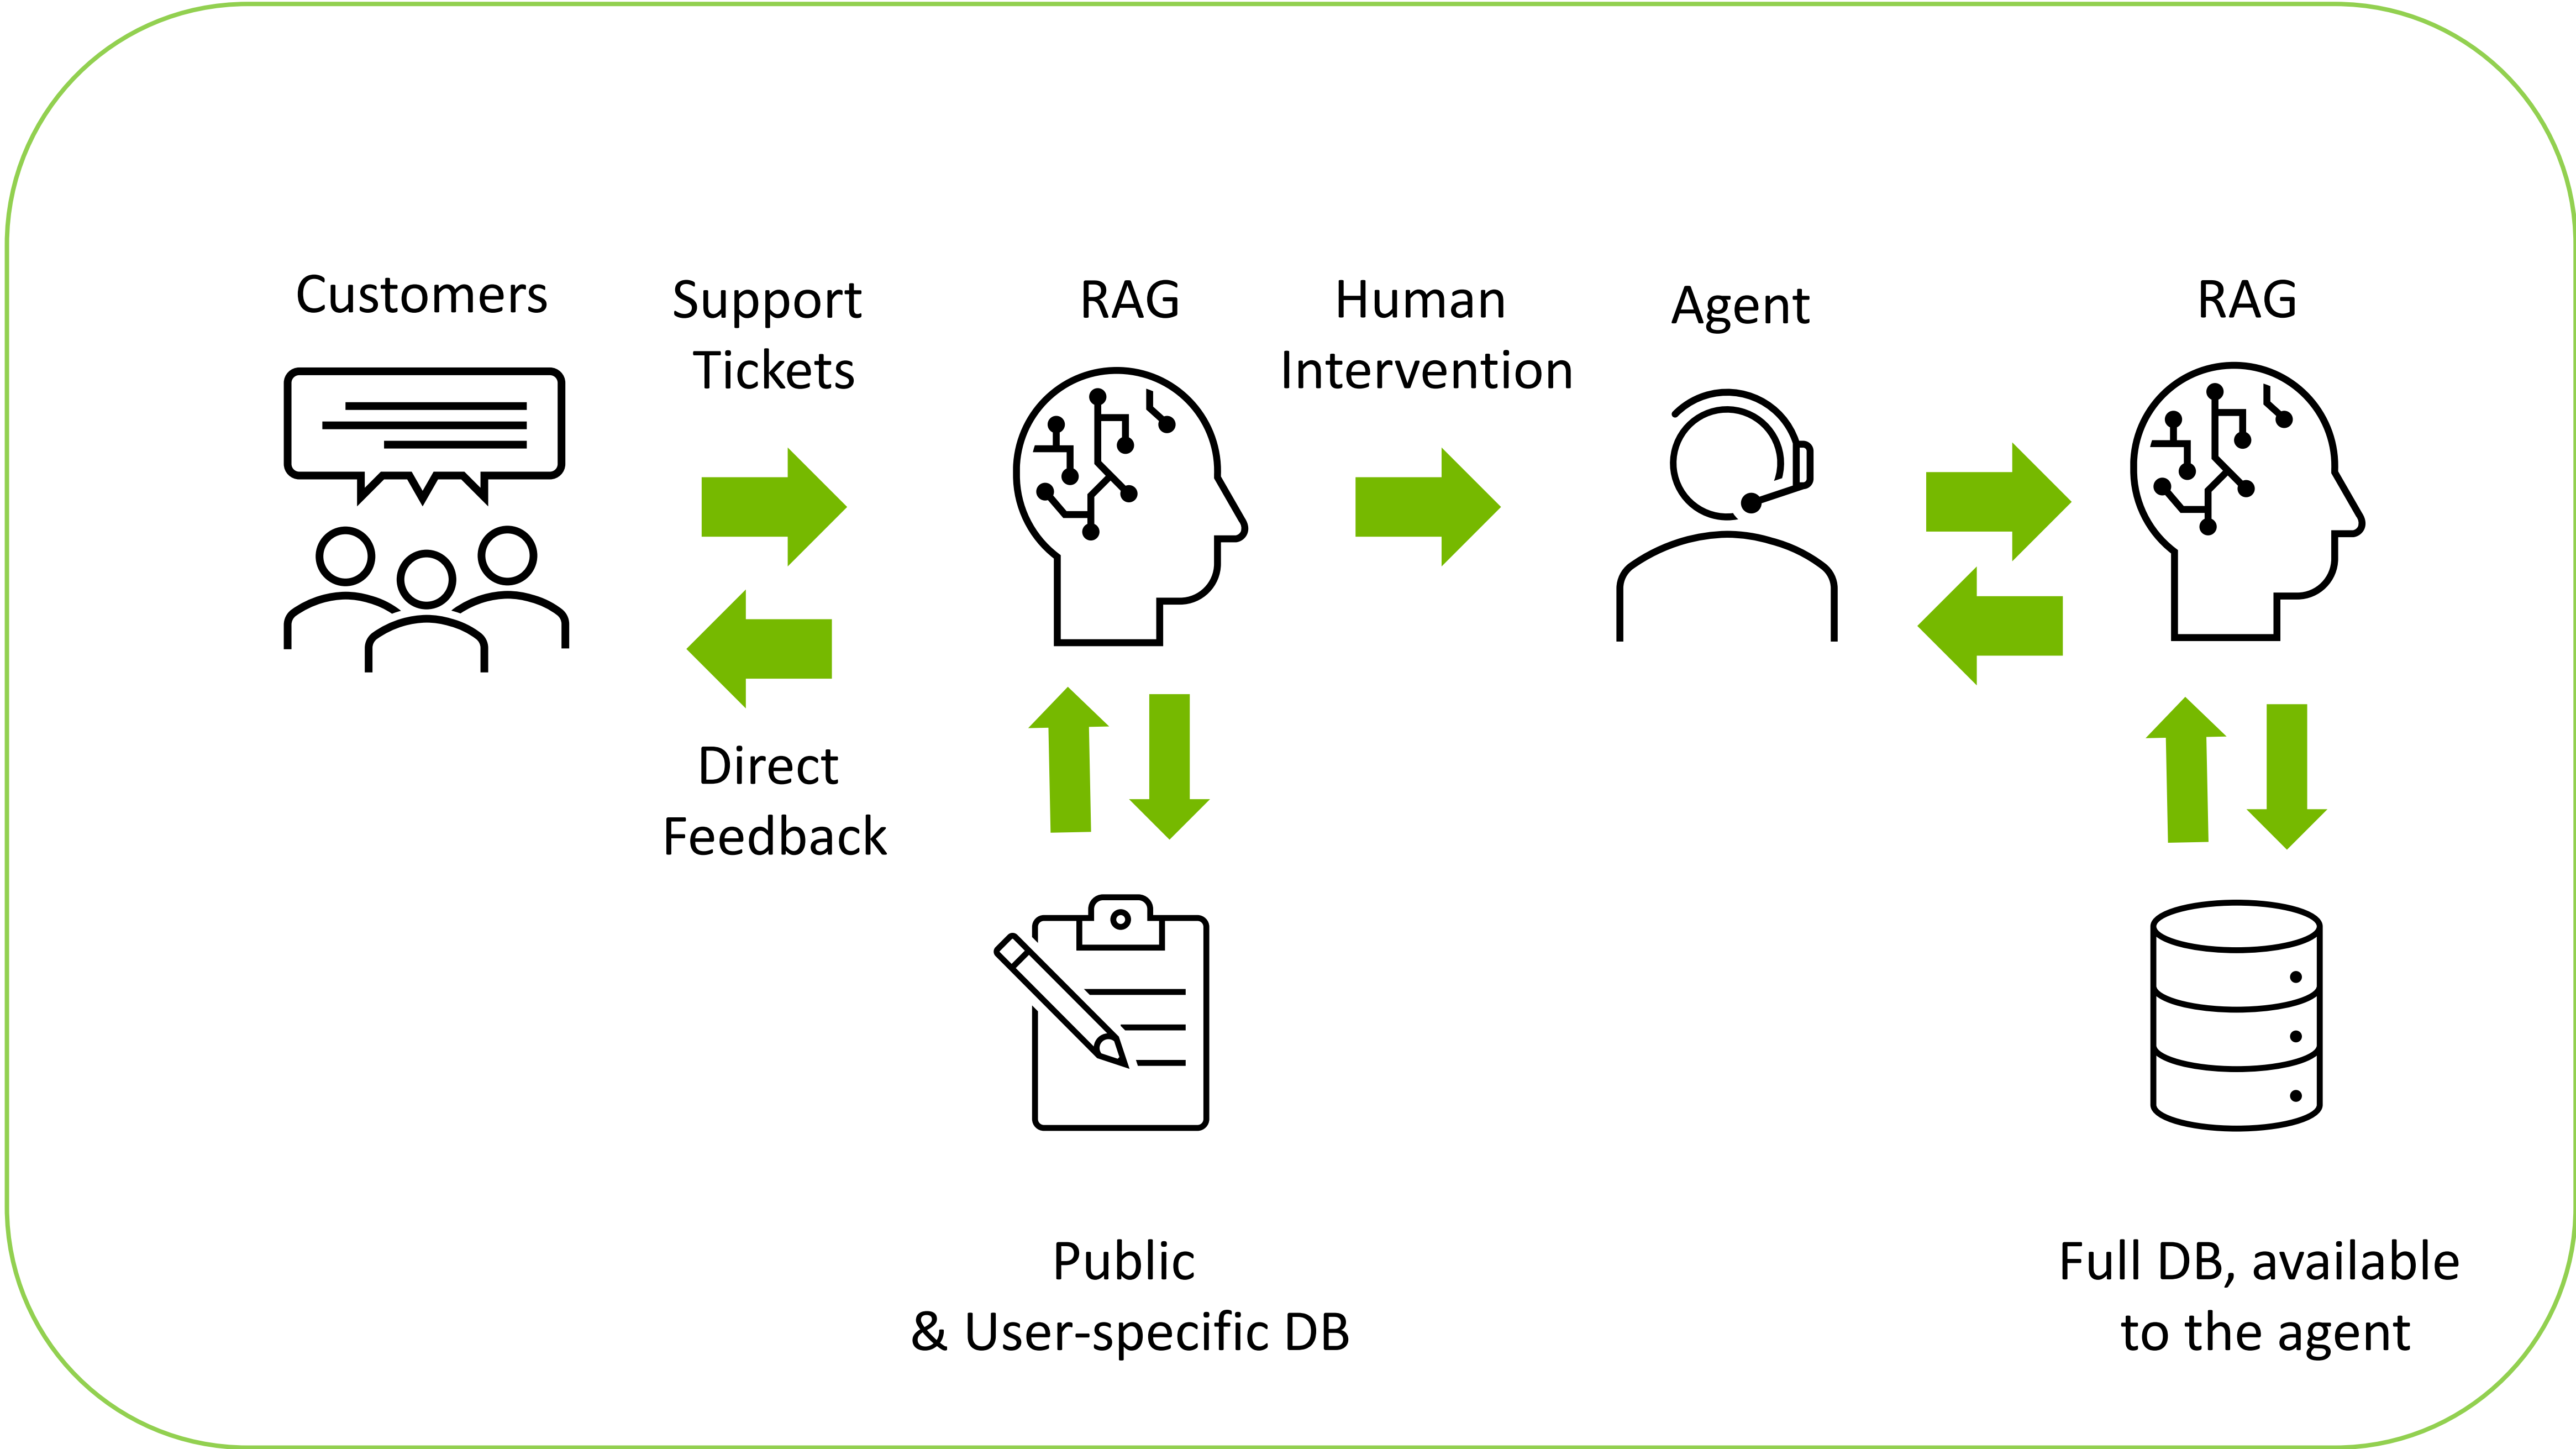
NVIDIA SW Stack for inference

Short summary of how to think about a problem



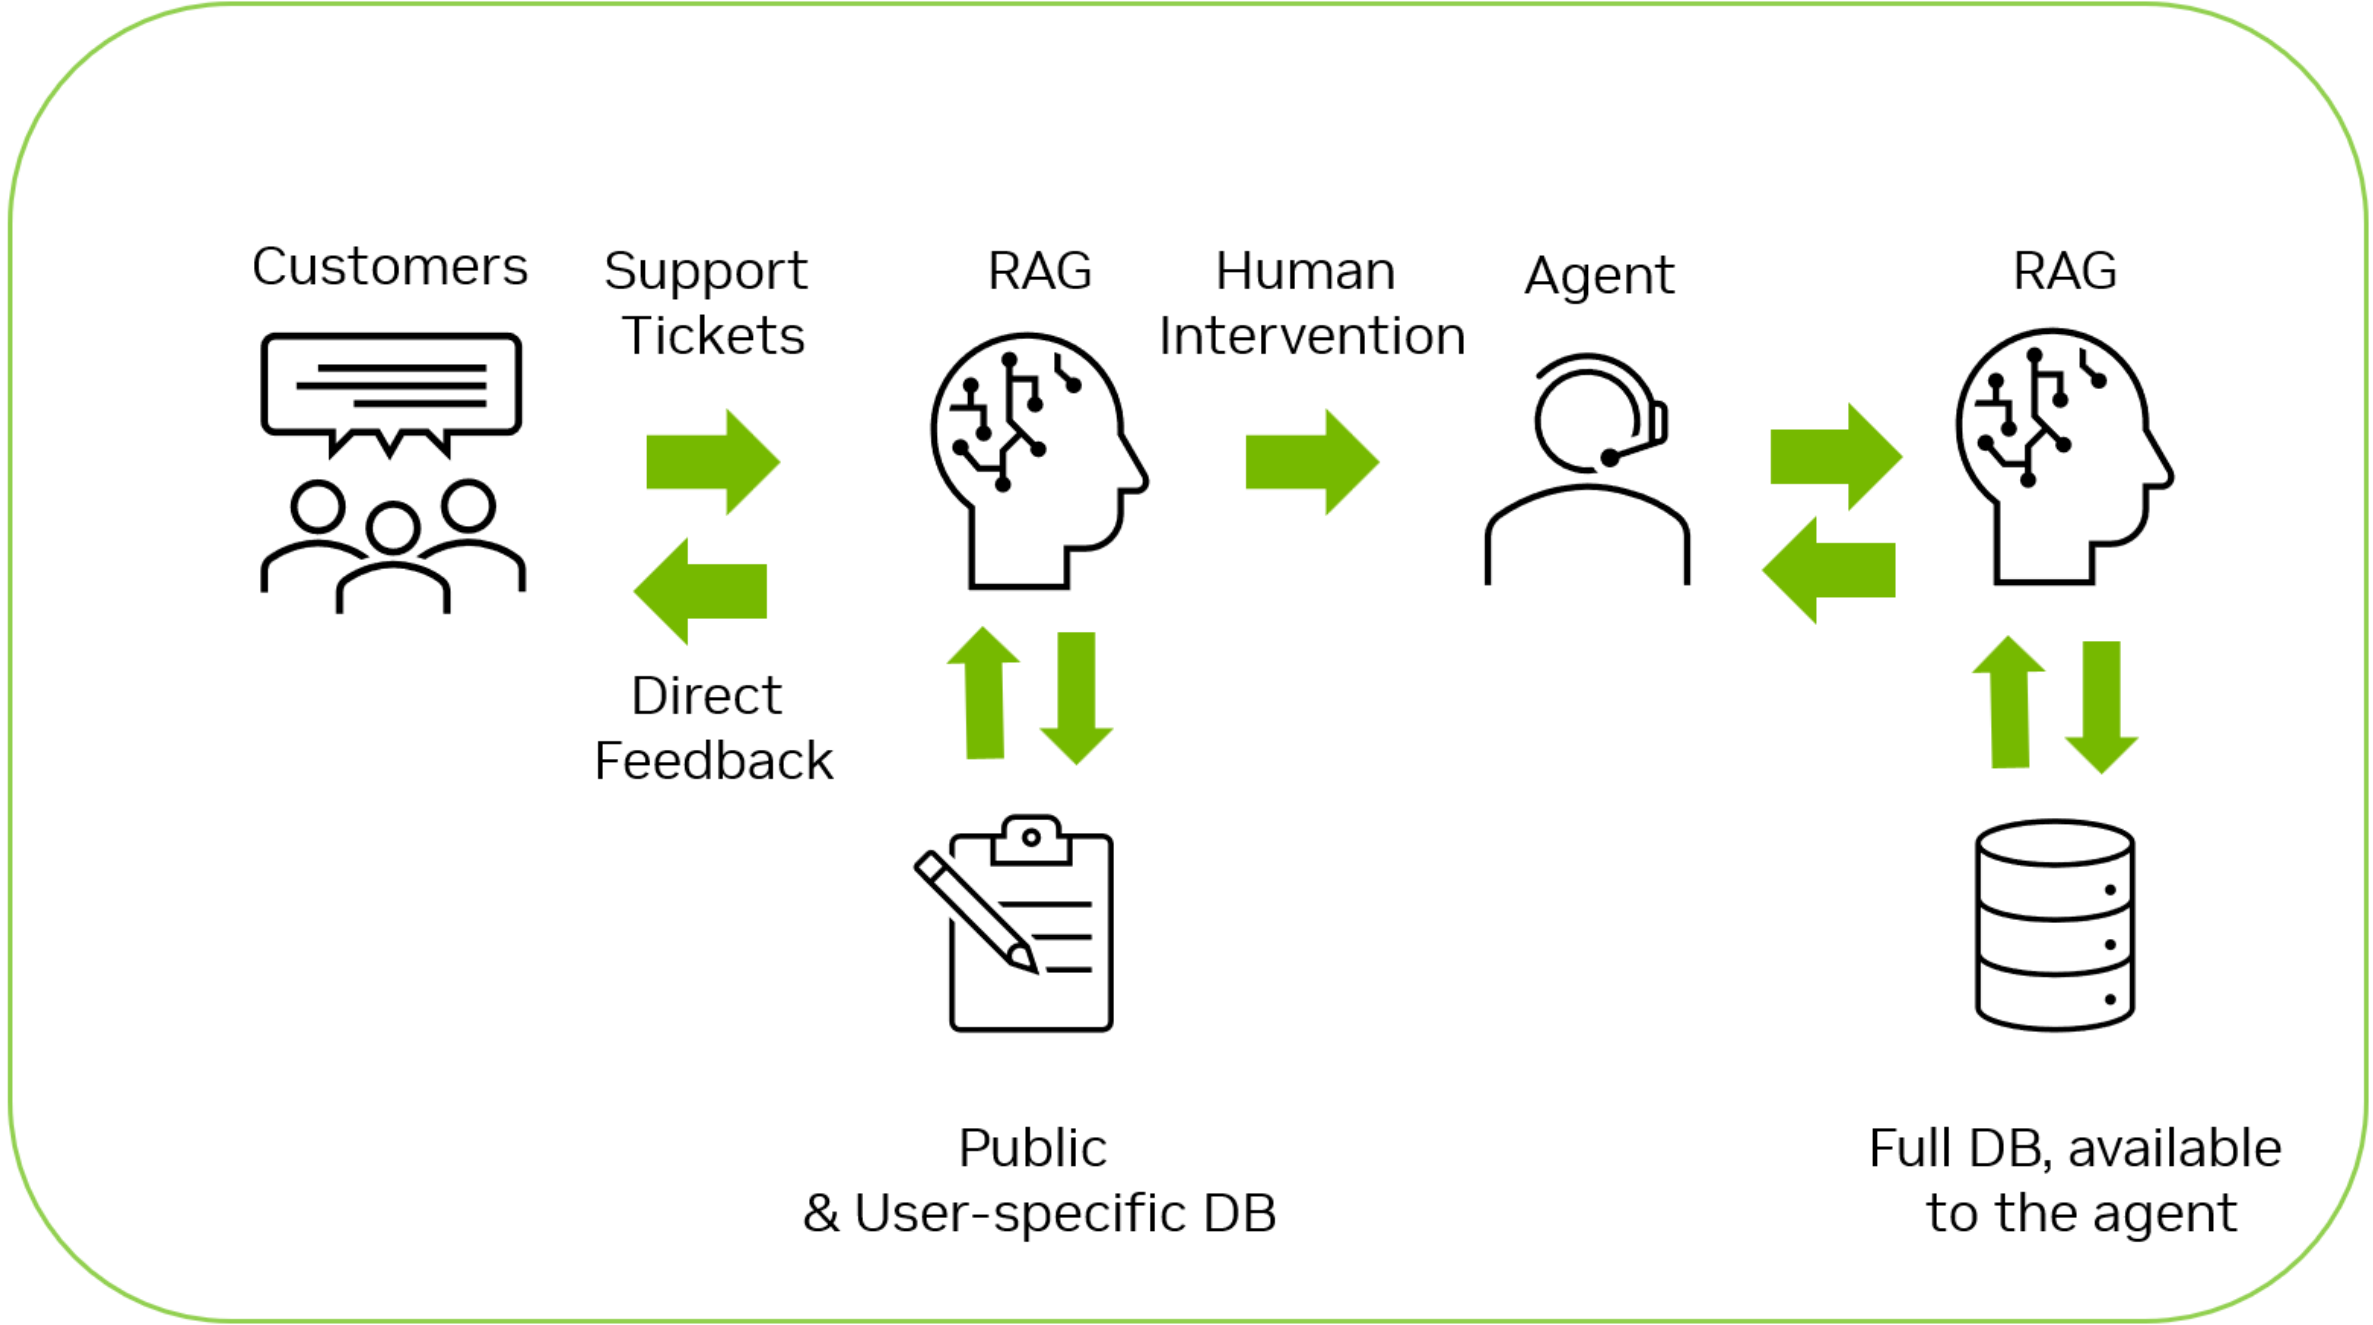
Customer Use Case Example

Challenges of sizing




Customer Use Case Example

Challenges of sizing




How many systems do we need to buy for this?



Gather requirements

- 3500 words in, 500 words out
- NeMo 43B GPT
- First token latency limit 3s
- Max 31 requests (=prompts) per second



Inference sizing

The customer needs 13 DGX H100 systems

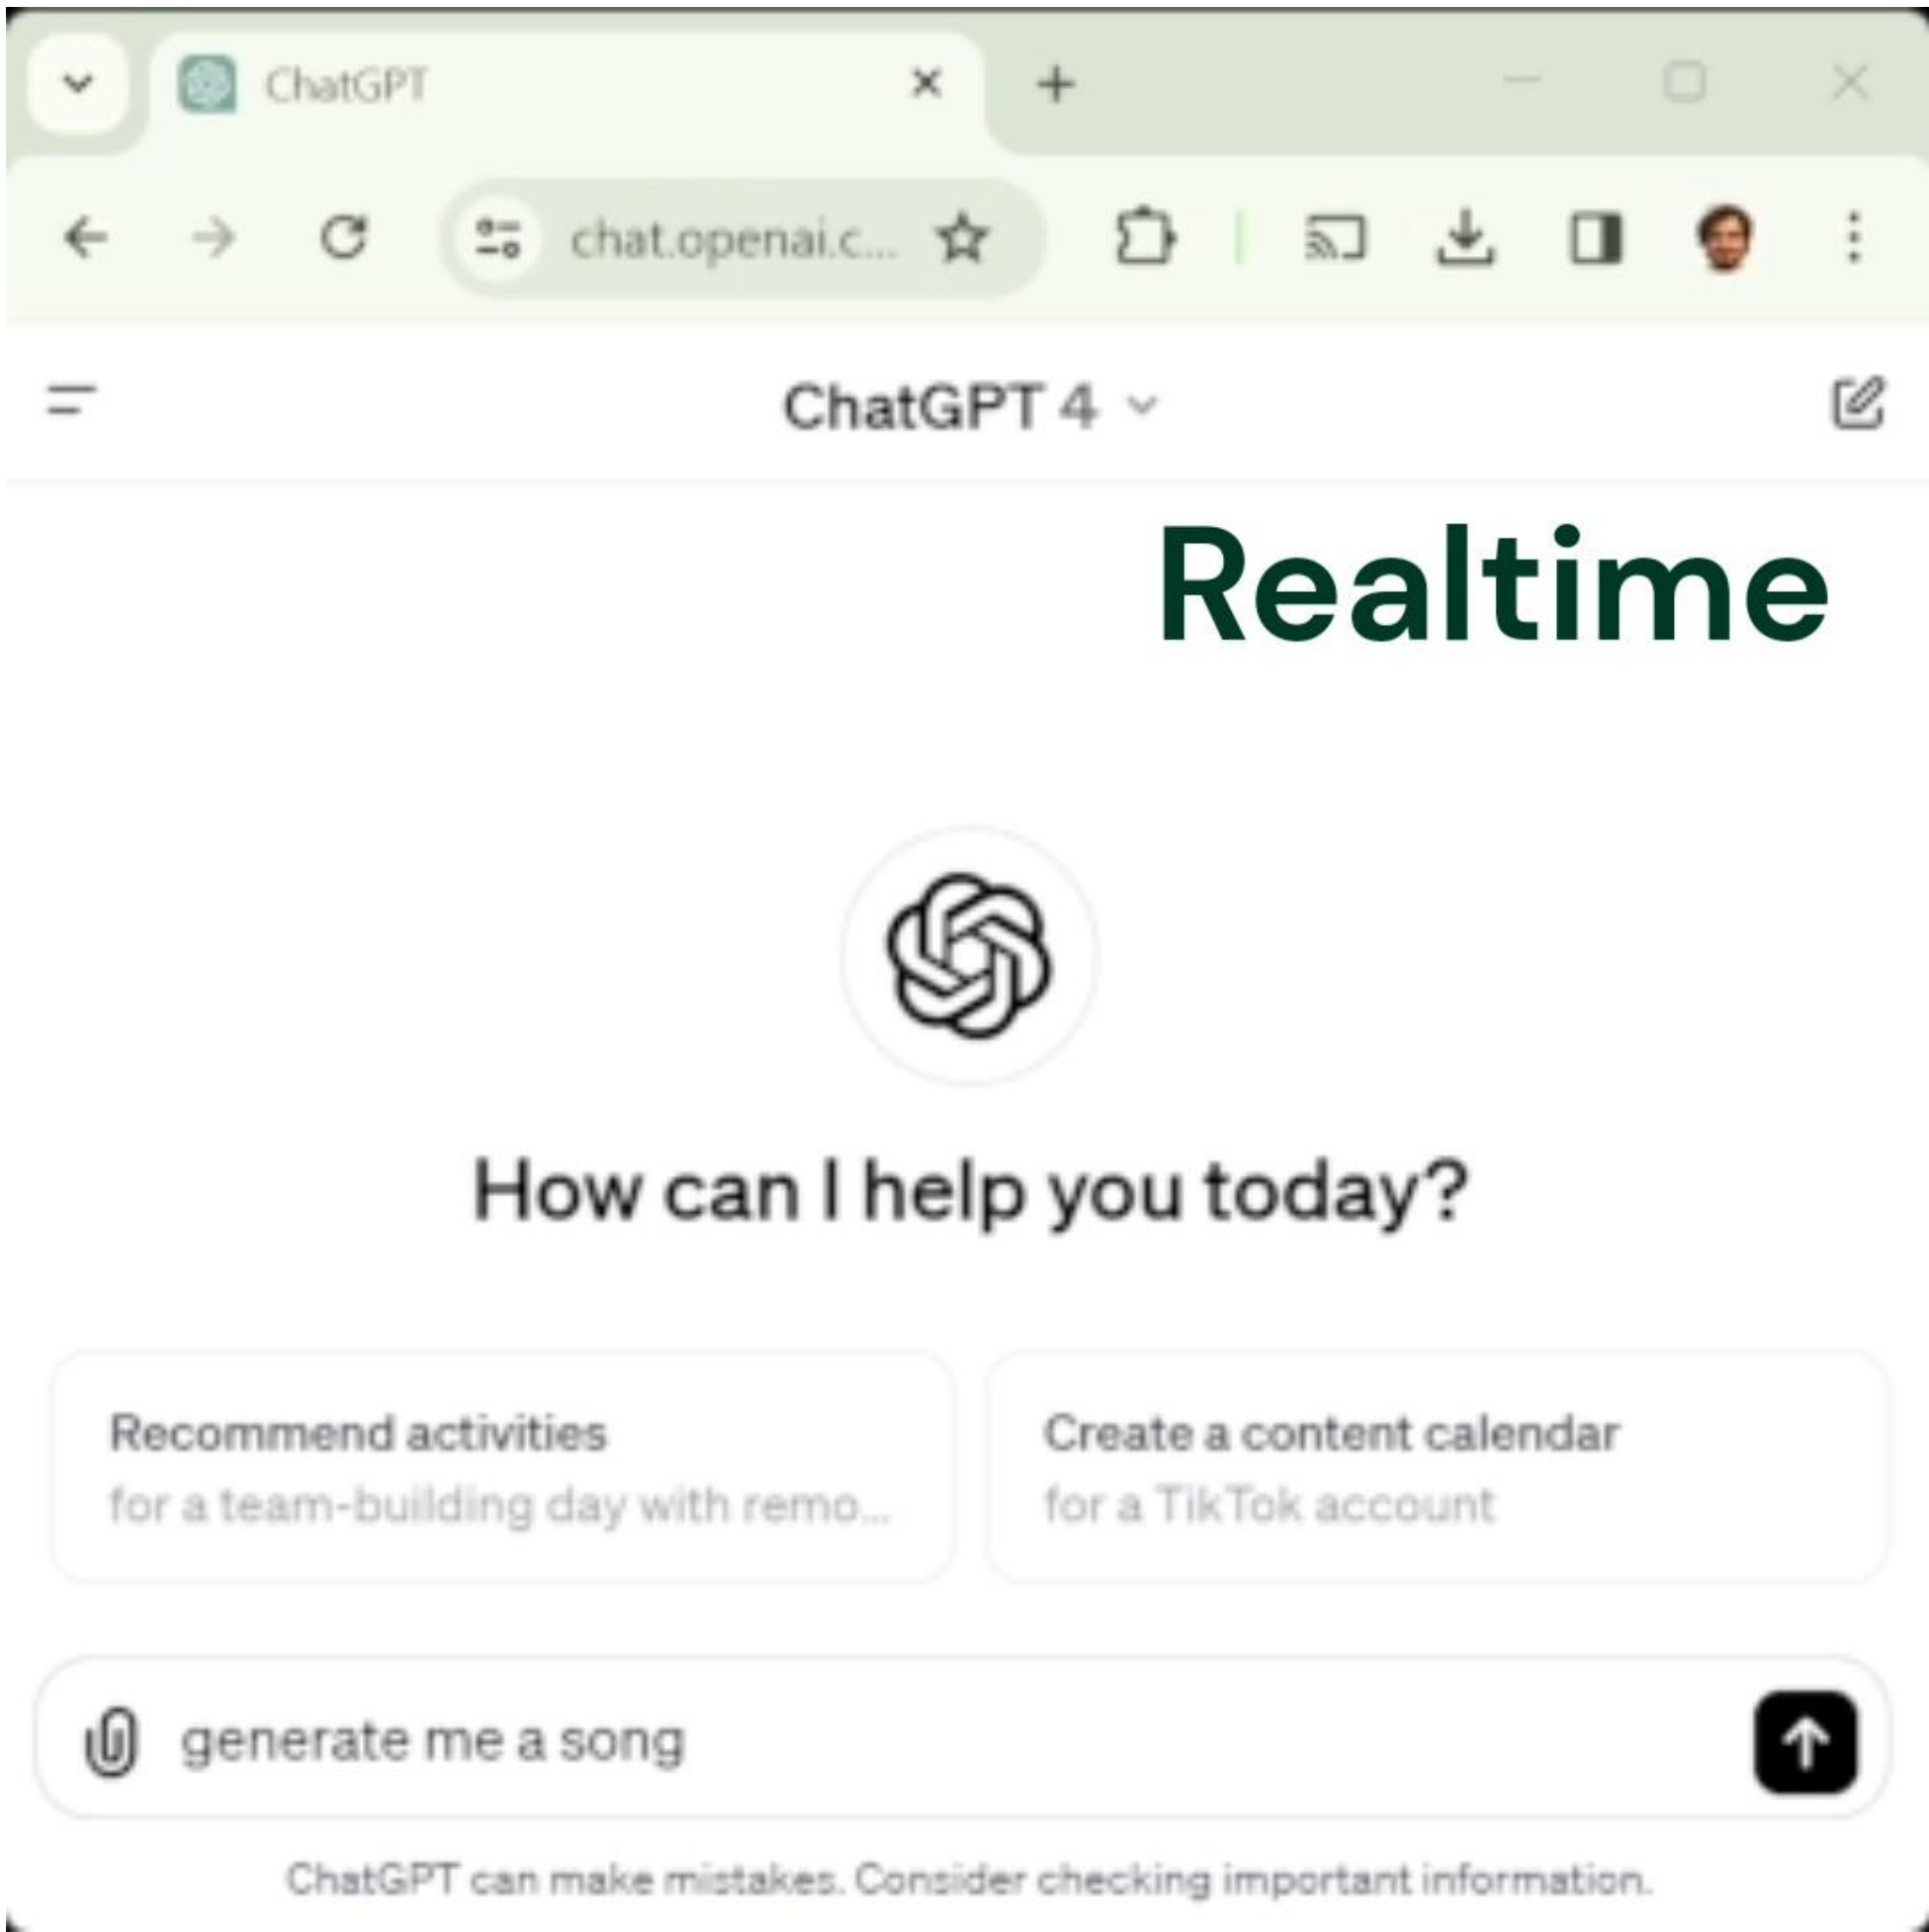
- Throughput: 2.4 requests per second
- First token latency 2606 ms (prefill) is within the limit specified
- Inter-token latency 21.4 ms/generated token
- Generation latency of 500 tokens = $21.4 * 500 = 10\,700\text{ ms} = 10.7\text{ s}$



Two Stages of LLM Execution

Prefill vs Decoding

- **Prefill** = time to first token (~word)
 - Loading the user prompt into the system
 - From the request reception to the first token
 - Depends only on the number of input tokens
 - Populate KV-cache for all the tokens from the prompt.
 - Compute-bound for most of the reasonable prompt lengths
- **Decoding** = inter-token latency
 - Generating the response token by token, word by word
 - Inter-token latency depends on the total token number, both input and output tokens.
 - Usually memory-bound




Prefill: 1.14s, 5 input tokens, 1 output token

generate me a song

Decoding: 1.62s, 33 output tokens

Could you please provide me with some sp


Prefill, first-response	Decoding							
User: generate me a song. AI: <i>Could</i>	you	please	generate	text	about	GPUs	that	




The Two Things To Care About

Where and how do we execute inference?

Where?



How?



Where?

Significant impact of deployment location


On-prem

- Fixed Capacity: you need to understand the size for the maximum simultaneous load
- Pricing model: per peak capacity

Cloud

- Variable Capacity: APIs hide capacity concerns – in reality, similar limitations apply (GPU shortage)
- Pricing model: per token

- Minimal Capacity + Autoscaling for bursts



How?

Significant impact of inference strategy

Online — live generation

- Complexity: it matters to people how quickly they will get their response
- Imposing latency requirement significantly decreases available throughput. Need to balance between throughput and latency

Offline — postponed computation

- Simplest execution model
- Throughput, throughput, throughput: maximum GPU utilization, maximum batch size

Fun fact: Fast human reading speed is 90 ms/token (=500 words/minute at 0.75 tokens/word) (avg is 200 ms/token)



Online Streaming vs Sequential

Two facets of latency

- **Streaming:** one token at a time
 - In this situation only the **TIME-TO-FIRST-TOKEN** matters (as we generate text faster than people can read)
 - One needs to develop the app streaming capabilities
 - Simpler to satisfy real-time latency requirements
 - Can be implemented only in the last step of the pipeline
- **Sequential:** waits for the full response
 - Say you want to check whether the user question is not toxic BEFORE you start answering
 - In this case **END-TO-END** latency/time to last token matters
 - Legacy apps can be simply updated with sequential mode
 - Latency requirements are too restricting for throughput




Questions to start sizing for inference



Questions for a Sizing Use Case


1. ✓ What model are you planning to use?
2. ✓ What is the average number of tokens in the prompt to your LLM (Length of input)?
 - For English one token is approximately 0.75 of a word.
 - Make sure to include system prompt.
3. ✓ What is the average number of tokens in your LLM output?
4. ✓ How many requests per second should your system process at its peak?
5. ✓ What is your latency limit? First-token? Last-token?
6. ✓ What GPUs are you considering?



Which Model?

The most popular requests

- Typically, we get asked about Llama 3 family of the models
 - Free for research and commercial use
 - Supported by NVIDIA SW stack, including NeMo, NIM, TRT-LLM and Triton
- The bigger the model, the more resources it needs for inference
 - The bigger the model the better the accuracy
 - Very roughly, the resource amount scales with the model size
- If considering Mistral 7B or Llama-3 8B parameters, see also NVIDIA Nemotron-3 8B Family of models: [blog](#)



Input Length

There's a maximum budget of tokens to pass into the model

SPO

- Most of the models support up to 4096 tokens.
 - Context window = input tokens + output tokens
 - Llama2 supports 4096 context window
 - New models support even larger context windows
- Everything counts so be careful:
 - **System prompt** (a.k.a custom instructions): instructions you give to the model for every “dialogue”. Make sure to include them into the input token count as shown in example on the right
 - **Retrieved documents** (a.k.a Retrieval Augmented Generation, RAG). For RAG pipelines key paragraphs from the internal document storage are added to the prompt, before the user requests. Typically RAG systems target to use full available context length
 - For 4K context typical 3500 input tokens, 500 output tokens
 - [What is RAG — NVIDIA blog](#)
 - **Chat history**: previous exchange of messages in the conversation

Custom instructions ⓘ

What would you like ChatGPT to know about you to provide better responses?

I work for NVIDIA as a Solutions Architect.

This system prompt costs +9 input tokens

43/1500

How would you like ChatGPT to respond?


Respond concisely, unless asked to expand your thoughts.

This sentence costs +12 input tokens

56/1500


Enable for new chats ☒

Cancel



Peak Requests Per Second

- Poisson distribution approximation
 - One knows the average, but would like to know the peak
- Find 95th percentile: [ChatGPT dialogue](#)



```
from scipy.stats import poisson
```

```
# Parameters
```

```
lambda_ = 64 # average number of requests per second
```

```
percentile = 0.95 # 95th percentile
```

```
# Calculate the 95th percentile value
```

```
k_95th_percentile = poisson.ppf(percentile, lambda_)
```

```
print(k_95th_percentile) # 77, 20% difference
```

```
print(poisson.ppf(0.95, 7)) # 12, 71% difference
```




LLM Inference Requires Multiple GPUs

Tensor Parallelism (TP) – so how to split your neural network

- Tensor Parallelism (TP) can be used for LLM Inference. One model gets split across several GPUs, which heavily relies on data exchange between GPUs
 - Lower latency, but lower throughput
 - $TP \geq 2$ required for bigger models like LLaMa-70B
- If $TP > 2$ we strongly recommend NVLink-enabled servers for inference, such as HGX and DGX systems (in contrast to PCIe servers)
- We normalize all the results for servers with 8 GPUs (even for L40S)
 - An instance is the group of GPUs forming a data replica of the model
 - $(\# \text{ of instances}) * TP = 8$
 - 8 instances with TP1, 2 instances with TP4

TP8	Instance 1							
TP4	In. 1				In. 2			
TP2	In. 1		In. 2		In. 3		In. 4	
TP1	In. 1	In. 2	In. 3	In. 4	In. 5	In. 6	In. 7	In. 8



Time = \$

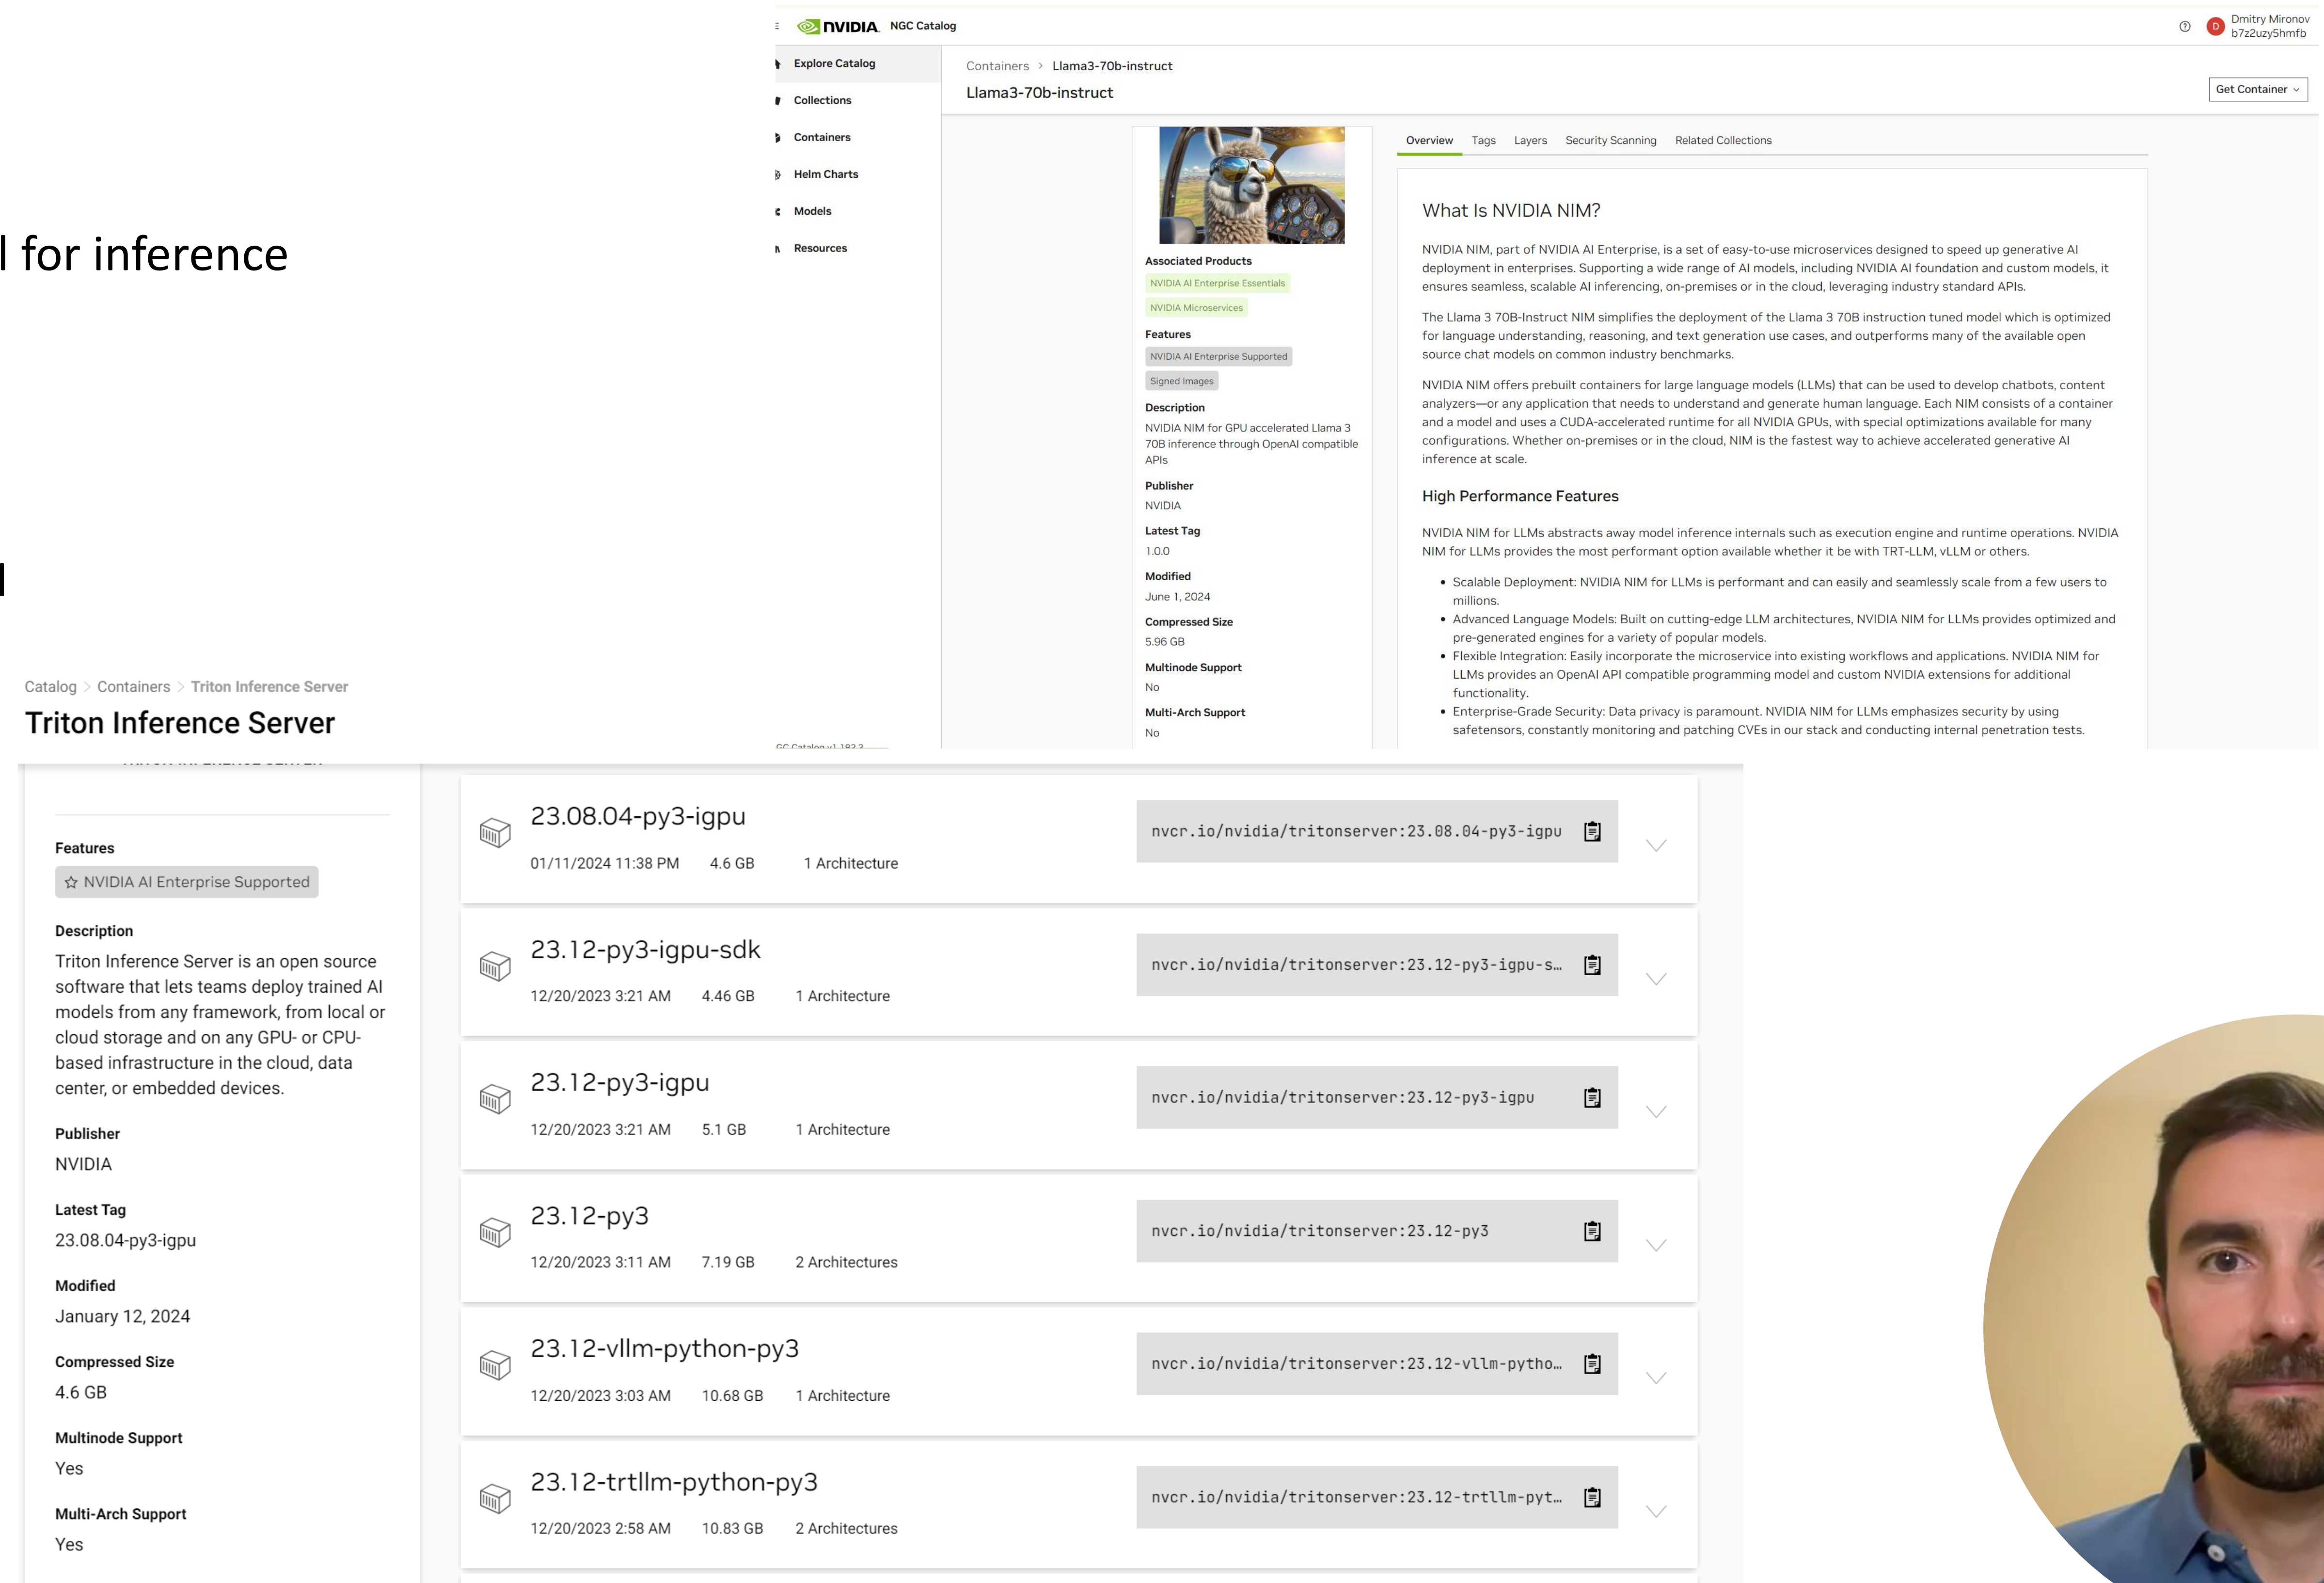


Tools Available

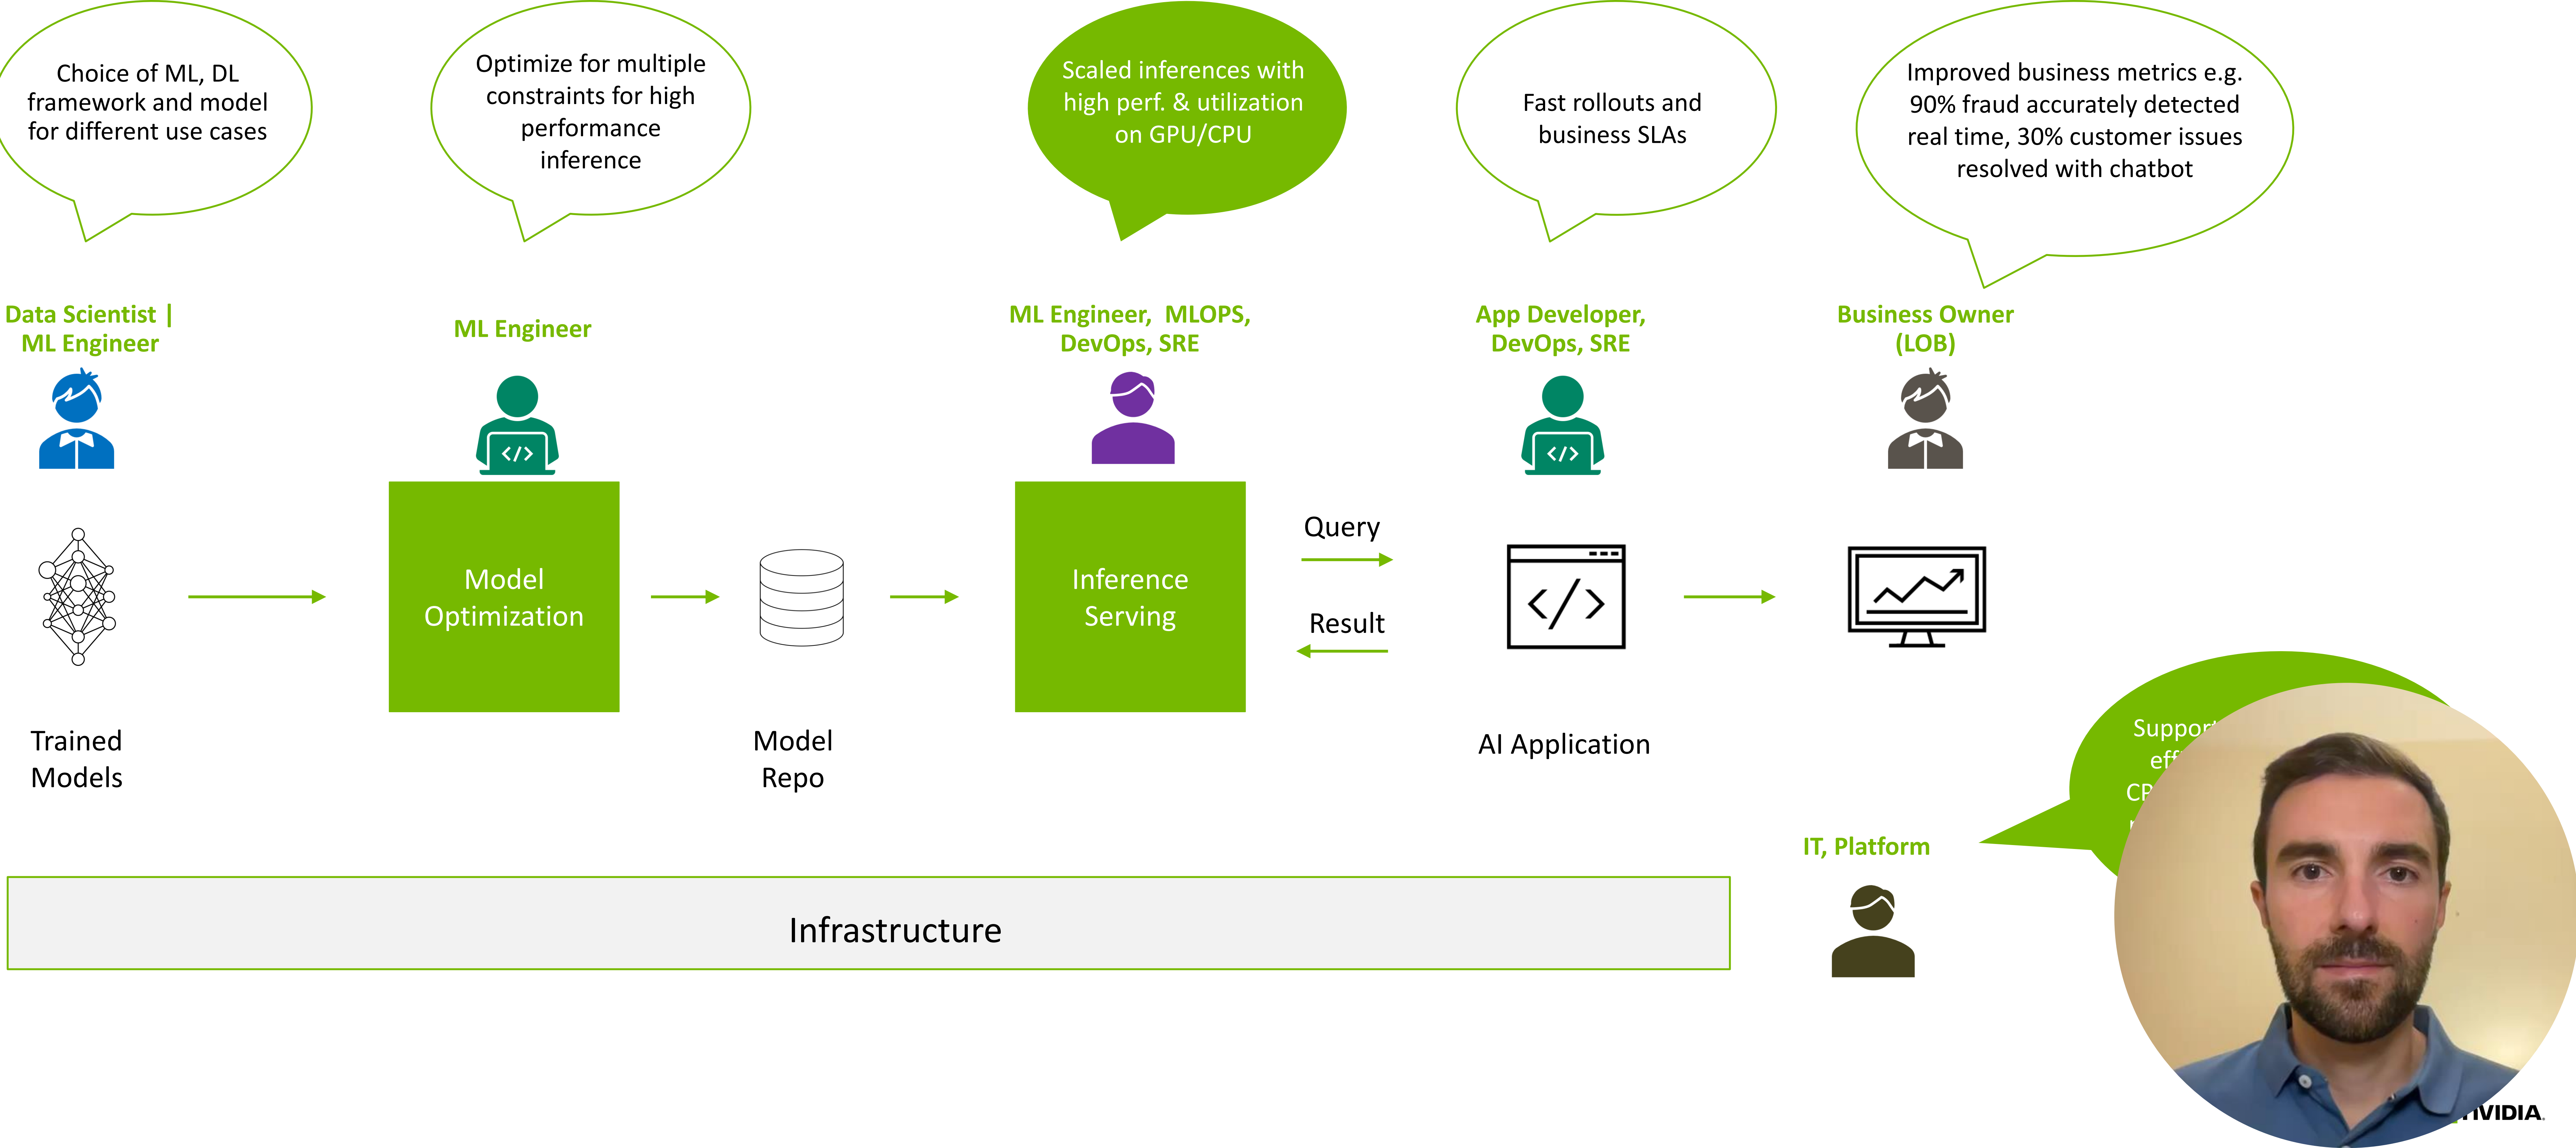


Inference Containers

- Triton + TensorRT-LLM
 - Open Source hands-on tools
 - TensorRT-LLM optimizes a model for inference
 - Triton is an inference server
- NVIDIA NIM
 - Deploy a LLM within minutes
 - Supports OpenAI-compatible API
 - Accelerated by TRT-LLM



AI Inference Workflow



TensorRT-LLM Optimizing LLM Inference


SoTA Performance for Large Language Models for Production Deployments

Challenges: LLM performance is crucial for real-time, cost-effective, production deployments. Rapid evolution in the LLM ecosystem, with new models & techniques released regularly, requires a performant, flexible solution to optimize models

TensorRT-LLM is an **open-source** library to **optimize inference performance** on the latest **Large Language Models** for NVIDIA GPUs. It is built on FasterTransformer and TensorRT with a simple Python API for defining, optimizing, & executing LLMs for inference in production

SoTA Performance

Leverage TensorRT compilation & kernels from FasterTransformer, CUTLASS, OAI Triton, ++



Ease Extension

Add new operators or models in Python to quickly support new LLMs with optimized performance

```
# define a new activation
def silu(input: Tensor) -> Tensor:
    return input * sigmoid(input)

#implement models like in DL FWs
class LlamaModel(Module)
    def __init__(...)
        self.layers = ModuleList([...])



    def forward (...)
        hidden = self.embedding(...)

        for layer in self.layers:
            hidden_states = layer(hidden)

        return hidden
```

LLM Batching with Triton

Maximize throughput and GPU utilization through new scheduling techniques for LLMs




NVIDIA

Triton Inference Server

Open-Source Software For Fast, Scalable, Simplified Inference Serving

Any Framework




Supports Multiple
Framework Backends
Natively e.g., TensorFlow,
PyTorch, TensorRT, XGBoost,
ONNX, Python & More

Any Query Type



Optimized for Real Time,
Batch, Streaming, Ensemble
Inferencing

Any Platform




X86 CPU | Arm CPU |
NVIDIA GPUs | MIG

Linux | Windows |
Virtualization

Public Cloud, Data Center
and Edge/Embedded
(Jetson)

DevOps & MLOps



Integration With
Kubernetes, KServe,
Prometheus & Grafana

Available Across All Major
Cloud AI Platforms

Performance & Utilization



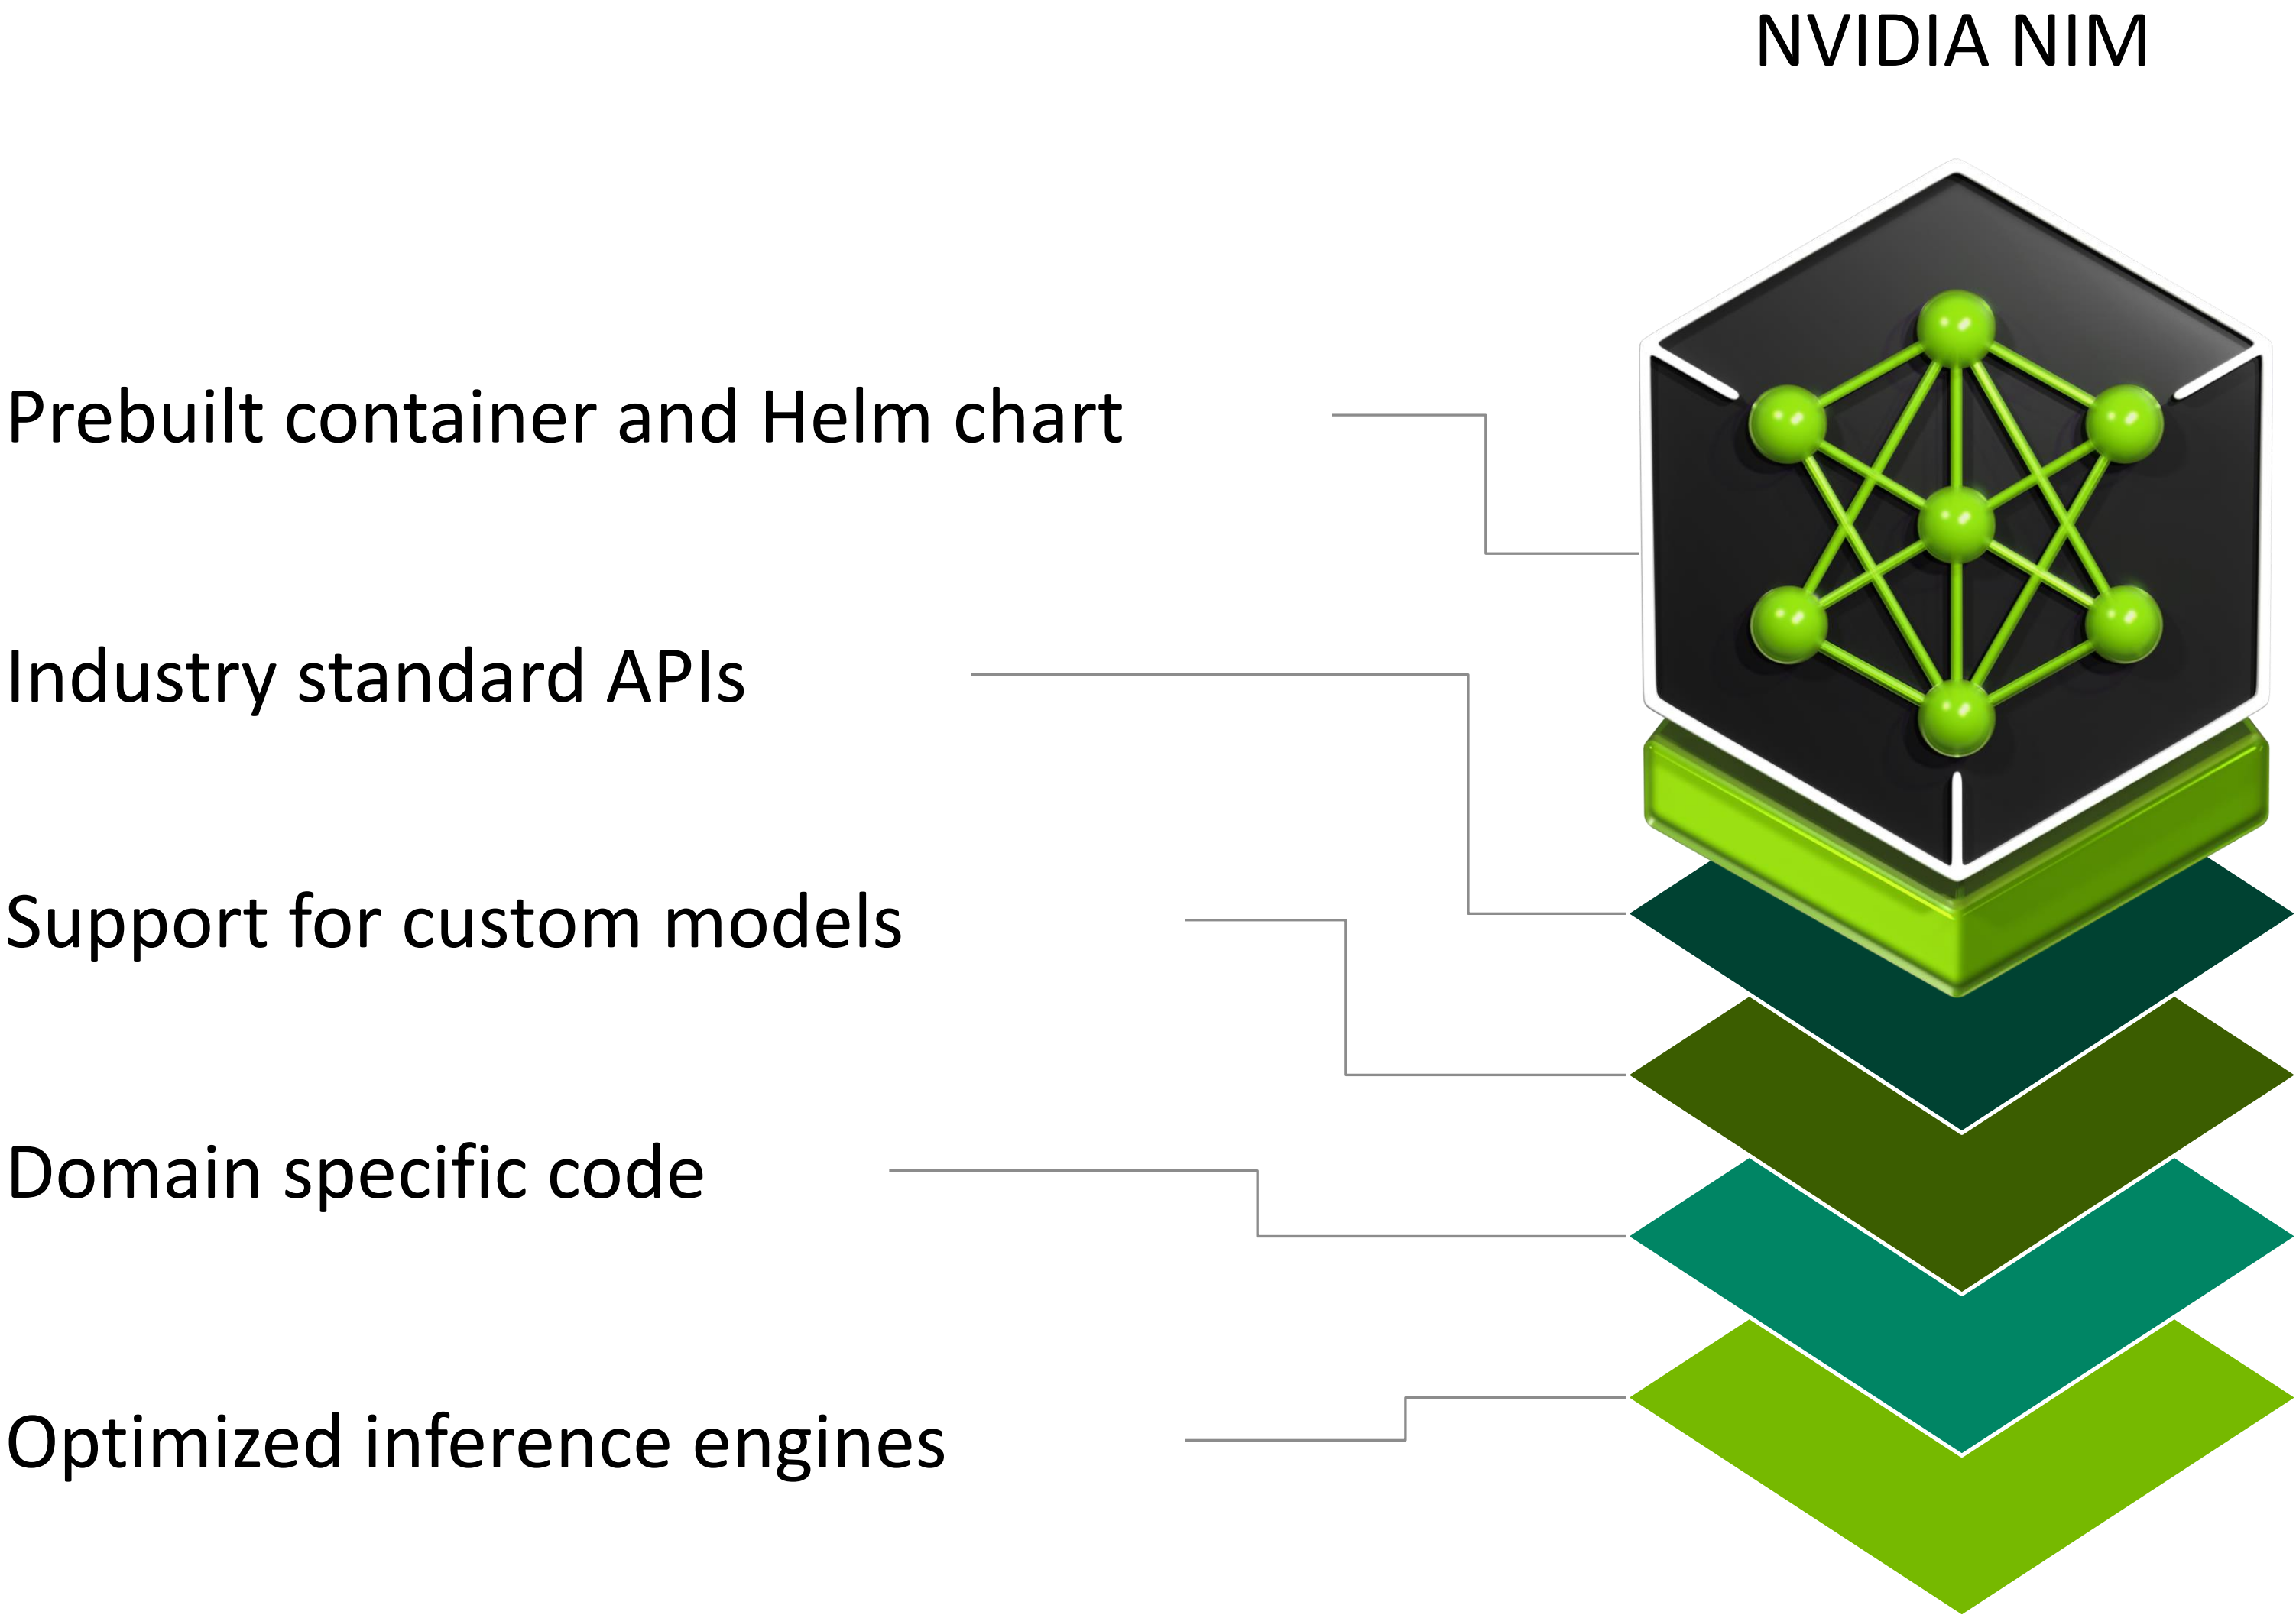
Model Analyzer for Optimal
Configuration

Optimal GPU/CPU
Throughput



NVIDIA NIM Optimized Inference Microservices

Accelerated runtime for generative AI



Deploy anywhere with security and control of AI applications and data

Speed time to market with prebuilt, continuously maintained microservices

Empower developers with the latest AI models, standard APIs and enterprise tools

Optimize throughput and latency to maximize token generation and responsiveness

Boost accuracy by tuning custom models from proprietary data sources

Deploy in production with API stability, quality assurance and enterprise support



DGX &
DGX Cloud



Publicly Available Performance Benchmarking

- Most recommended: GenAI-Perf from Triton team
 - https://github.com/triton-inference-server/client/tree/main/src/c%2B%2B/perf_analyzer/genai-perf
 - Triton GenAI-Perf is a CLI tool which can help you optimize the inference performance of models running on Triton Inference Server and OpenAI endpoints by measuring changes in performance as you experiment with different optimization strategies.
 - Used in NIM for LLMs Performance Guide <https://docs.nvidia.com/nim/benchmarking/llm/latest/index.html>
- <https://github.com/NVIDIA/TensorRT-LLM/tree/main/benchmarks/cpp> — TensorRT-LLM C++
 - TensorRT-LLM provides users with an easy-to-use Python API to define Large Language Models (LLMs) and build TensorRT engines that contain state-of-the-art optimizations to perform inference efficiently on NVIDIA GPUs. TensorRT-LLM also contains components to create Python and C++ runtimes that execute those TensorRT engines.
 - Some results: <https://github.com/NVIDIA/TensorRT-LLM/blob/main/docs/source/performance.md>
- Triton CLI for limited experimentation: https://github.com/triton-inference-server/triton_cli



Example of benchmarking




Example with Llama 3 8B

Smaller model – for auxiliary task

- We are looking for a sizable use case of **Llama3-8B**. 2000 in, 200 out, **TTFT < 500ms**
- For input 2000, output 200 we have **70.7** peak prompts per second per **one** DGX H100
- That's 2M requests per working day (8 hours)
- 3 requests per person → **679k daily active users**
- 4.2B input, 261M output tokens per day

meta-llama3-8b-instruct, H100_80GB_HBM3, input length: 2000, output length: 200



Best performance per optimal scenario

Metric	Throughput-optimal ▼	Latency-optimal ▲
First Token Latency (ms)	489.5	29.9

Prompts per Second per 8 GPUs	70.7	3.7
Latency per Generated Token (ITL, TPOT) (ms)	11.7	5.24
Prompts per Second per 1 GPU	8.8	0.47
Cost Per 1000 Prompts (USD)	0.053	1
Cost Per 1M Input Tokens (USD)	0.021	0.39
Cost Per 1M Output Tokens (USD)	0.062	1.2



Rules of Thumb for Sizing

- We estimate the sizing based on NVIDIA SW stack: NIM or TensorRT-LLM (=TRT-LLM) and Triton Inference Server
- For models greater than 13B, that need more than 1 GPU, prefer NVLink-enabled systems
- In the streaming mode, when the words are returned one by one, first-token latency is determined by the input length
- The cost and the latency are usually dominated by the number of output tokens
 - Example below: H100 SXM, Llama 70B, BS 8, TP 4, FP 16.
Input of 3500 tokens takes the same amount of time as generating 99 tokens
(2.6 seconds each stage, 26.8 ms/generated token)
 - However, generating is almost always faster than human reading speed
 - Thus, input tokens are much cheaper
- Introducing latency limit can significantly decrease available throughput
- Larger models require more memory and have higher latency, scaling approximately with the model size

Input processing: 3500 tokens

Generating 99 tokens out



Inference Resources

- NIM for LLM Benchmarking Guide <https://docs.nvidia.com/nim/benchmarking/llm/latest/index.html>
- NVIDIA NIM: <https://docs.nvidia.com/nim/index.html>
- GTC session about LLM inference sizing: <https://www.nvidia.com/en-us/on-demand/session/gtc24-s62797/>
- Mastering LLM Techniques: Inference Optimization— NVIDIA Blog <https://developer.nvidia.com/blog/mastering-llm-techniques-inference-optimization/>



